

Artificial Intelligence

Machine Learning

Ensemble Methods



Outline

1. Majority voting
2. Boosting
3. Bagging
4. Random Forests
5. Conclusion

Majority Voting

- A randomly chosen hyperplane has an expected **error of 0.5**.

Majority Voting

- A randomly chosen hyperplane has an expected **error of 0.5**.
- Many random hyperplanes combined by **majority vote** will still be random.

Majority Voting

- A randomly chosen hyperplane has an expected **error of 0.5**.
- Many random hyperplanes combined by **majority vote** will still be random.
- Suppose we have m classifiers, performing slightly better than random, that is **error = $0.5 - \epsilon$** .

Majority Voting


- A randomly chosen hyperplane has an expected **error of 0.5**.
- Many random hyperplanes combined by **majority vote** will still be random.
- Suppose we have m classifiers, performing slightly better than random, that is **error = $0.5 - \epsilon$** .
- Combine: make a decision based on majority vote?

Majority Voting

- A randomly chosen hyperplane has an expected **error of 0.5**.
- Many random hyperplanes combined by **majority vote** will still be random.
- Suppose we have m classifiers, performing slightly better than random, that is **error = $0.5 - \epsilon$** .
- Combine: make a decision based on majority vote?
- What if we combined these m **slightly-better-than-random** classifiers? Would majority vote be a good choice?

Condorcet's Jury Theorem

Marquis de Condorcet Application of Analysis to the Probability of Majority Decisions. 1785.




Assumptions:

1. Each individual makes the right choice with a probability p .
2. The votes are independent.

Condorcet's Jury Theorem

Marquis de Condorcet Application of Analysis to the Probability of Majority Decisions. 1785.



Assumptions:

1. Each individual makes the right choice with a probability p .
2. The votes are independent.

If $p > 0.5$, then adding more voters increases the probability that the majority decision is correct. if $p < 0.5$, then adding more voters makes things worse.

Ensemble Methods

- An **Ensemble Method** combines the predictions of many individual classifiers by majority voting.

Ensemble Methods

- An **Ensemble Method** combines the predictions of many individual classifiers by majority voting.
- Such individual classifiers, called **weak learners**, are required to perform slightly better than random.

Ensemble Methods

- An **Ensemble Method** combines the predictions of many individual classifiers by majority voting.
- Such individual classifiers, called **weak learners**, are required to perform slightly better than random.

**How do we produce independent weak learners
using the same training data?**

Ensemble Methods

- An **Ensemble Method** combines the predictions of many individual classifiers by majority voting.
- Such individual classifiers, called **weak learners**, are required to perform slightly better than random.

**How do we produce independent weak learners
using the same training data?**

- Use a **strategy** to obtain relatively independent weak learners!
- Different methods:
 1. Boosting
 2. Bagging
 3. Random Forests

Boosting

- First ensemble method.
- One of the most powerful Machine Learning methods.
- Popular algorithm: AdaBoost.M1 (Freund and Shapire 1997).
- “Best off-the-shelf classifier in the world” Breiman (CART’s inventor), 1998.
- Simple algorithm.
- Weak learners can be trees, perceptrons, decision stumps, etc.
- **Idea:**

Train the weak learners on weighted training examples.

Boosting

Boosting

- The predictions from all of the G_m , $m \in \{1, \dots, M\}$ are combined with a weighted majority voting.
- α_m is the contribution of each weak learner G_m .
- Computed by the boosting algorithm to give a weighted importance to the classifiers in the sequence.
- The decision of a highly-performing classifier in the sequence should weight more than less important classifiers in the sequence.
- This is captured in:

$$G(x) = \text{sign}\left[\sum_{m=1}^M \alpha_m G_m(x)\right]$$

Boosting

The error rate on the training sample:

$$err := \frac{\sum_{i=1}^n 1\{y_i \neq G(x_i)\}}{n}$$

Boosting

The error rate on the training sample:

$$err := \frac{\sum_{i=1}^n 1\{y_i \neq G(x_i)\}}{n}$$

The error rate on each weak learner:

$$err_m := \frac{\sum_{i=1}^n w_i 1\{y_i \neq G_m(x_i)\}}{\sum_{i=1}^n w_i}$$

Boosting

The error rate on the training sample:

$$err := \frac{\sum_{i=1}^n 1\{y_i \neq G(x_i)\}}{n}$$

The error rate on each weak learner:


$$err_m := \frac{\sum_{i=1}^n w_i 1\{y_i \neq G_m(x_i)\}}{\sum_{i=1}^n w_i}$$

Intuition:

- Give large weights for hard examples.
- Give small weights for easy examples.

Boosting

For each weak learner m , we associate an error err_m .



$$\alpha_m = \log\left(\frac{1 - err_m}{err_m}\right)$$

AdaBoost

1. Initialize the example weights $w_i = \frac{1}{n}$, $i = 1, \dots, n$.

AdaBoost

1. Initialize the example weights $w_i = \frac{1}{n}$, $i = 1, \dots, n$.
2. For $m = 1$ to M (number of weak learners)

AdaBoost

1. Initialize the example weights $w_i = \frac{1}{n}$, $i = 1, \dots, n$.
2. For $m = 1$ to M (number of weak learners)
 - (a) Fit a classifier $G_m(x)$ to training data using the weights w_i .

AdaBoost

1. Initialize the example weights $w_i = \frac{1}{n}$, $i = 1, \dots, n$.
2. For $m = 1$ to M (number of weak learners)
 - (a) Fit a classifier $G_m(x)$ to training data using the weights w_i .
 - (b) Compute

$$err_m := \frac{\sum_{i=1}^n w_i \mathbf{1}\{y_i \neq G_m(x_i)\}}{\sum_{i=1}^n w_i}$$

AdaBoost

1. Initialize the example weights $w_i = \frac{1}{n}$, $i = 1, \dots, n$.
2. For $m = 1$ to M (number of weak learners)
 - (a) Fit a classifier $G_m(x)$ to training data using the weights w_i .
 - (b) Compute

$$err_m := \frac{\sum_{i=1}^n w_i \mathbf{1}\{y_i \neq G_m(x_i)\}}{\sum_{i=1}^n w_i}$$

- (c) Compute

$$\alpha_m = \log\left(\frac{1 - err_m}{err_m}\right)$$

AdaBoost

1. Initialize the example weights $w_i = \frac{1}{n}$, $i = 1, \dots, n$.
2. For $m = 1$ to M (number of weak learners)
 - (a) Fit a classifier $G_m(x)$ to training data using the weights w_i .
 - (b) Compute

$$err_m := \frac{\sum_{i=1}^n w_i \cdot 1\{y_i \neq G_m(x_i)\}}{\sum_{i=1}^n w_i}$$

- (c) Compute

$$\alpha_m = \log\left(\frac{1 - err_m}{err_m}\right)$$

- (d) $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot 1(y_i \neq G_m(x_i))]$ for $i = 1, \dots, n$.

AdaBoost

1. Initialize the example weights $w_i = \frac{1}{n}$, $i = 1, \dots, n$.
2. For $m = 1$ to M (number of weak learners)
 - (a) Fit a classifier $G_m(x)$ to training data using the weights w_i .
 - (b) Compute

$$err_m := \frac{\sum_{i=1}^n w_i \mathbf{1}\{y_i \neq G_m(x_i)\}}{\sum_{i=1}^n w_i}$$

- (c) Compute

$$\alpha_m = \log\left(\frac{1 - err_m}{err_m}\right)$$

- (d) $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot \mathbf{1}(y_i \neq G_m(x_i))]$ for $i = 1, \dots, n$.

3. Output

$$G(x) = \text{sign}\left[\sum_{m=1}^M \alpha_m G_m(x)\right]$$

Digression: Decision Stumps

This is an example of very weak classifier


A simple 2-terminal node decision tree for binary classification.

$$f(x_i|j, t) = \begin{cases} +1 & \text{if } x_{ij} > t \\ -1 & \text{otherwise} \end{cases}$$

Where $j \in \{1, \dots, d\}$.

Example: A dataset with 10 features, 2,000 examples training and 10,000 testing.


AdaBoost Performance



Error rates:


- Random: 50%.
- Stump: 45.8%.
- Large classification tree: 24.7%.
- AdaBoost with stumps: 5.8% after 400 iterations!

AdaBoost Performance



AdaBoost with Decision stumps lead to a form of:
feature selection

AdaBoost-Decision Stumps



Bagging & Bootstrapping

- Bootstrap is a re-sampling technique \equiv sampling from the empirical distribution.
- Aims to improve the quality of estimators.
- Bagging and Boosting are based on bootstrapping.
- Both use re-sampling to generate weak learners for classification.
- **Strategy:** Randomly distort data by re-sampling.
- Train weak learners on re-sampled training sets.
- Bootstrap **aggregation** \equiv **Bagging**.

Bagging

Training

For $b = 1, \dots, B$

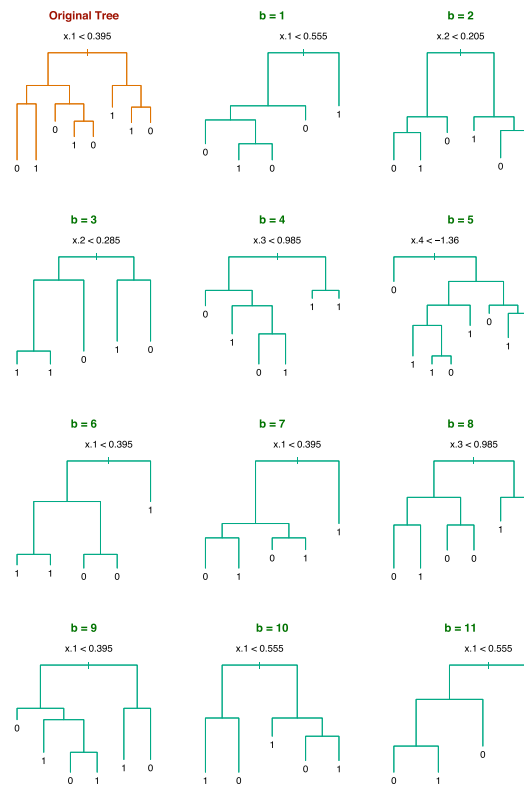
1. Draw a bootstrap sample \mathcal{B}_b of size ℓ from training data.
2. Train a classifier f_b on \mathcal{B}_b .

Classification: Classify by majority vote among the B trees:

$$f_{avg} := \frac{1}{B} \sum_{b=1}^B f_b(x)$$

Bagging

Bagging works well for trees:



Random Forests

1. Random forests: modifies bagging with trees to reduce correlation between trees.
2. Tree training optimizes each split over all dimensions.
3. But for Random forests, **choose a different subset of dimensions at each split**. Number of dimensions chosen m .
4. Optimal split is chosen within the subset.
5. The subset is chosen at random out of all dimensions $1, \dots, d$.
6. Recommended $m = \sqrt{d}$ or smaller.

Credit

- “An Introduction to Statistical Learning, with applications in R” (Springer, 2013)” with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.